



MODELING VNFS IN A CONTRAIL ENVIRONMENT

How to design a VNF using Contrail & Heat

Umberto Manfredini

JUNIPER
NETWORKS

Engineering
Simplicity

AGENDA

- Where are we now?
- What's Heat?
- Contrail and Heat: modeling VNFs
- From vanilla Neutron to Juniper Contrail
- Service Chaining




JUNIPER CONTRAIL

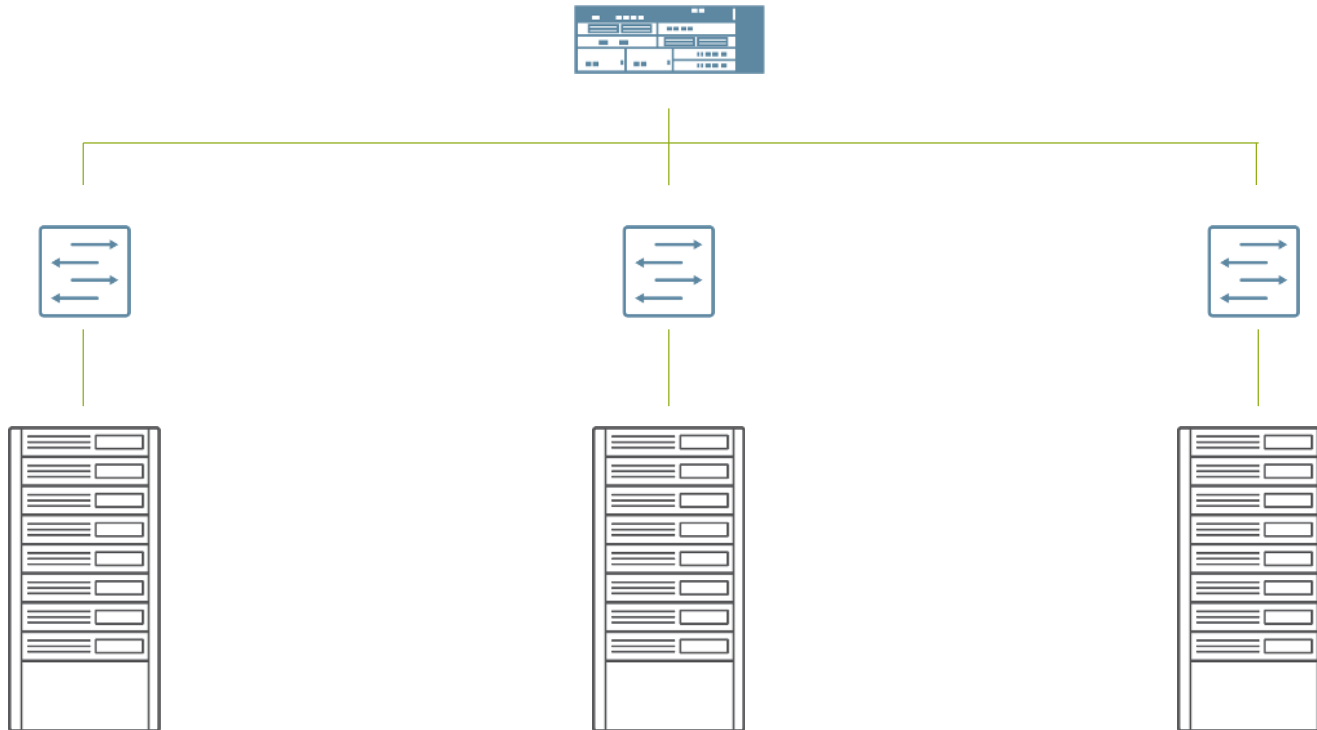
- Juniper SDN Controller
- Manages virtual networking within the DC
- Based on well known protocols and solutions like BGP and VPNs



ORIGINALLY IT WAS A SWITCHED NETWORK




THEN DATA CENTERS CAME



FINALLY VIRTUALIZATION APPEARED

- A network within a network
- A network inside a server
- Virtual Machines are the new appliances
- Virtual Networks are the new switched networks



SAME THING...DIFFERENT OBJECTS




CALL FOR NEW MODELS

- Yesterday, we used cables to connect physical devices
- Now, we have to connect virtual machines
- We need new models to describe these new scenarios




VIRTUALIZATION BRINGS AUTOMATION

- Virtualization is software
- Software is more inclined to be automated
- Virtualization opens to network automation



AUTOMATION MEANS HEAT

- Openstack deploys and manage the DC
- Openstack has tool for automation called Heat
- Contrail fully integrates with Heat
- Contrail provides its own Heat objects



AND HEAT MEANS YAML

- Heat is a built-in Openstack module
- Heat uses YAML language
- Easy to read and write



YamL

CONTRAIL HAS ITS OWN OBJECTS

ubuntu [®]	admin ▼		
		OS::ContrailV2::ServiceTemplate	OpenStack ContrailV2
		OS::ContrailV2::Subnet	OpenStack ContrailV2
		OS::ContrailV2::VirtualDns	OpenStack ContrailV2
		OS::ContrailV2::VirtualDnsRecord	OpenStack ContrailV2
		OS::ContrailV2::VirtualIpp	OpenStack ContrailV2
		OS::ContrailV2::VirtualMachine	OpenStack ContrailV2
		OS::ContrailV2::VirtualMachineInterface	OpenStack ContrailV2
		OS::ContrailV2::VirtualNetwork	OpenStack ContrailV2
		OS::ContrailV2::VirtualRouter	OpenStack ContrailV2

AUTOMATION THROUGH STACKS

- Heat creates stacks
- Stacks are a list resources
- Resources are virtual network, virtual machines, etc...




WHAT DO I HAVE TO DO?



FROM WORDS TO HEAT

- Need to take an abstract input as request
- For example “configure a virtual network”
- And translate it into a Heat template




WHAT'S A HEAT TEMPLATE?

- Structured YAML files
- Template file, where we define resources
- Environment file, where we assign values to parameters




OBJECTS FOR EVERYONE

- Openstack offers a wide range of Heat object
- Nova objects
- Glance objects
- Contrail objects
- ...



INSIDE A HEAT TEMPLATE

- Heat templates made of multiple sections
- Two main sections
- Parameters section
- Resource section



DECLARING PARAMETERS

```
heat_template_version: 2015-04-30
```

```
parameters:
```

```
  myname :  
    type: string  
    default: 'namex'
```

```
  myimage :  
    type: string  
    default: 'ubuntu'
```

```
  myflavor :  
    type: string  
    default: 'smallvm'
```



DECLARING RESOURCES

resources:

test-net:

type: OS::Neutron::Net	→	resource type
properties:		
name: 'heat-net'	→	hardocded value

test-subnet:

type: OS::Neutron::Subnet		
properties:		
name: 'heat-subnet'		
network_id: { get_resource: test-net }	→	get_resource
cidr: '192.168.1.0/16'		


server1:

type: OS::Nova::Server		
properties:		
name: { get_param: myname }		
image: { get_param: myimage }	→	get_param
flavor: { get_param: myflavor }		
networks:		
- network: { get_resource: test-net }		

ENVIRONMENT FILE

- Parameters defined in the template file
- We need to give them a value
- Environment file is for this!


```
parameters:  
  myname : Ubuntu-1  
  myimg  : ubu14-iso  
  myflavor : linux-perf
```



WHO'S FIRST?

- Heat is smart
- But not too smart
- Some order is required
- Use “depends_on” to manage it

```
server1:
  type: OS::Nova::Server
  depends_on [test-subnet, test-net]
  properties:
    name: { get_param: myname }
    image: { get_param: myimage }
    flavor: { get_param: myflavor }
    networks:
      - network: { get_resource: test-net }
```




- Virtual network
- Virtual machine interface
- Instance IP
- BGPaaS
- Sub-interfaces
- Allowed address pair



CONTRAIL VIRTUAL NETWORKS

- Key element in Contrail networking
- Used to connect Virtual Machines



VIRTUAL NETWORK PROPERTIES

- Virtual Networks has a name
- Reference an IPAM (Contrail wrapper for virtual networks)

```
resources:  
  vip_vn:  
    type: OS::ContrailV2::VirtualNetwork  
    properties:  
      name: 'myvn'  
      network_ipam_refs: ['default-domain:default-project:default-network-ipam']
```



CUSTOMIZING THE VIRTUAL NETWORK

- Forwarding mode
- RPF check
- Flooding unknown unicast
- VXLAN identifier

resources:

vip_vn:

type: OS::ContrailV2::VirtualNetwork

properties:

virtual_network_properties:

{


virtual_network_properties_forwarding_mode: '13',

virtual_network_properties_rpf: 'enable',

virtual_network_properties_vxlan_network_identifier: '1234',

}


flood_unknown_unicast: 'true'



HOW DO I ECMP?


- Contrail supports ECMP
- Up to 5 configurable fields

```
resources:  
  vip_vn:  
    type: OS::ContrailV2::VirtualNetwork  
    properties:  
      ecmp_hashing_include_fields:  
        {  
          ecmp_hashing_include_fields_destination_ip: 'true',  
          ecmp_hashing_include_fields_destination_port: 'false',  
          ecmp_hashing_include_fields_hashing_configured: 'true',  
          ecmp_hashing_include_fields_ip_protocol: 'true',  
          ecmp_hashing_include_fields_source_ip: 'false',  
          ecmp_hashing_include_fields_source_port: 'false',  
        }  
      port_security_enabled: 'true'
```



CREATING SUBNETS

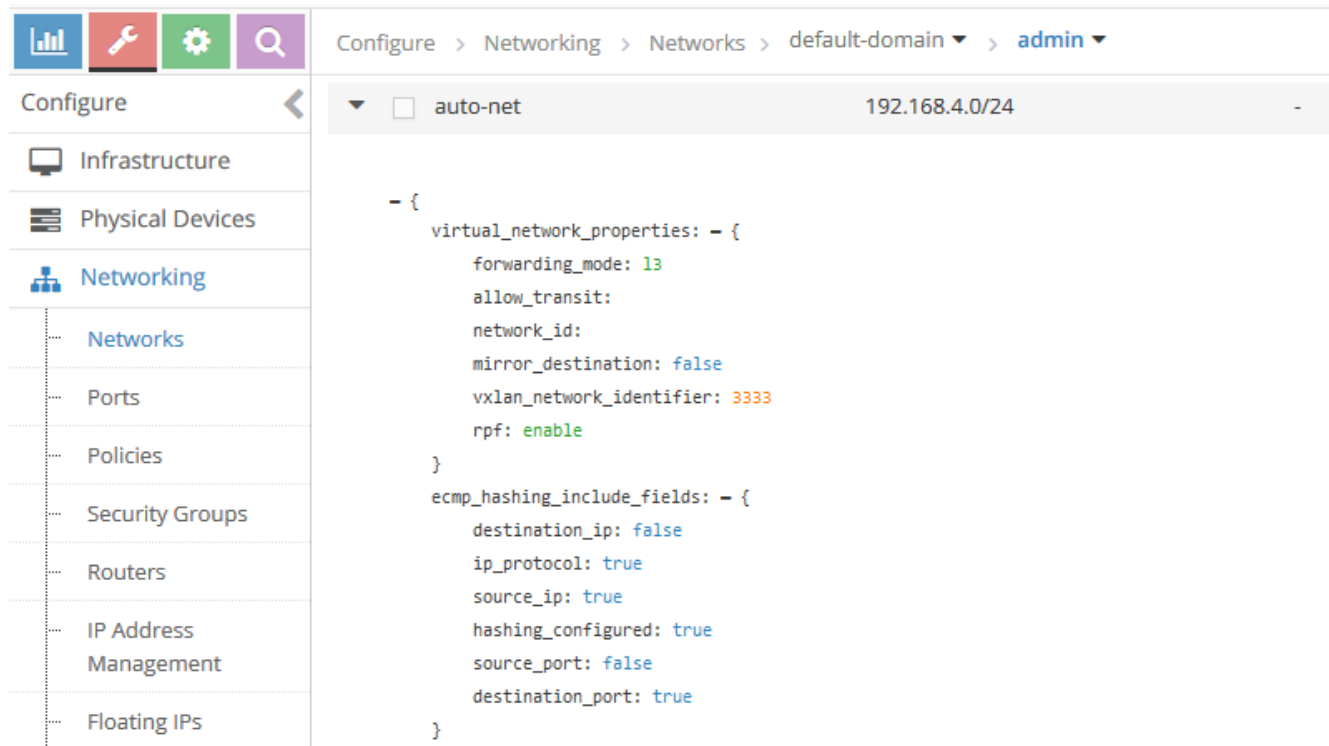
- Not a dedicated object
- Network address and mask
- Allocation pools
- DHCP
- Gateway
- Allocation order



SUBNETS IN HEAT

```
resources:
  vip_vn:
    type: OS::ContrailV2::VirtualNetwork
    properties:
      network_ipam_refs_data:
        [{
          network_ipam_refs_data_ipam_subnets:
            [{
              network_ipam_refs_data_ipam_subnets_subnet:
                {
                  network_ipam_refs_data_ipam_subnets_subnet_ip_prefix: '192.168.33.0',
                  network_ipam_refs_data_ipam_subnets_subnet_ip_prefix_len: '24',
                },
                network_ipam_refs_data_ipam_subnets_addr_from_start: 'true',
                network_ipam_refs_data_ipam_subnets_default_gateway: '192.168.33.3',
                network_ipam_refs_data_ipam_subnets_enable_dhcp: 'false',
                network_ipam_refs_data_ipam_subnets_allocation_pools:
                  [{
                    network_ipam_refs_data_ipam_subnets_allocation_pools_start: '192.168.33.5',
                    network_ipam_refs_data_ipam_subnets_allocation_pools_end: '192.168.33.10',
                  }]
            }]
        }]
    ]]
```


CHECKING EVERYTHING IS ALRIGHT



The screenshot displays the Juniper Cloud Managed Network (CMN) configuration interface. The left sidebar shows the navigation menu with the following items: Configure, Infrastructure, Physical Devices, Networking (selected), Networks, Ports, Policies, Security Groups, Routers, IP Address Management, and Floating IPs. The main content area shows the configuration path: Configure > Networking > Networks > default-domain > admin. Below this, a table lists the network configuration for 'auto-net' with the IP address 192.168.4.0/24. The configuration details are shown in a code block format.

Network	IP Address	Configuration
auto-net	192.168.4.0/24	-

```
- {
  virtual_network_properties: - {
    forwarding_mode: 13
    allow_transit:
    network_id:
    mirror_destination: false
    vxlan_network_identifier: 3333
    rpf: enable
  }
  ecmp_hashing_include_fields: - {
    destination_ip: false
    ip_protocol: true
    source_ip: true
    hashing_configured: true
    source_port: false
    destination_port: true
  }
}
```


VIRTUAL NIC

- Virtual machine port
- Connects VM to the Virtual Network
- It is actually 2 objects
- Virtual Machine Interface
- Instance IP



VIRTUAL MACHINE INTERFACE

- Minimal configuration
- Name
- Virtual Network Reference



```
resources:
  vmi_newnet:
    type: OS::ContrailV2::VirtualMachineInterface
    properties:
      name: 'my_vmi'
      virtual_network_refs: [{ list_join: [':', { get_attr: [ newnet_vn, fq_name ] } ] } ] }
```



Internal function
Fully Qualified Name
prepends "domain:project"

LEVERAGING EXISTING NETWORKS

- Existing network
- Simply pass the FQDN network name as a parameter
- FQDN name means “domain:project:network”

```
vmi_prex:  
  type: OS::ContrailV2::VirtualMachineInterface  
  properties:  
    virtual_network_refs: [ {get_param: mgmt-net} ]
```

```
parameters:  
  mgmt-net: 'default-domain:admin:mgmt'
```

env
file



I WANT MY OWN MAC

- By default, Contrail automatically assigns a MAC to a port
- Optionally, we can assign a custom MAC

```
resources:  
  vmi_newnet:  
    type: OS::ContrailV2::VirtualMachineInterface  
    properties:  
      virtual_machine_interface_mac_addresses:  
        {  
          virtual_machine_interface_mac_addresses_mac_address: [ 'aa:bb:cc:dd:ee:ff' ],  
        }
```



A PORT NEEDS AN IP

- Second element is the Instance IP
- References a virtual networks
- References a VMI

```
instanceip_1:
  type: OS::ContrailV2::InstanceIp
  depends_on: [ vmi_newnet, newnet_vn ]
  properties:
    virtual_machine_interface_refs: [{ get_resource: vmi_newnet }]
    virtual_network_refs: [{ list_join: [':', { get_attr: [ newnet_vn, fq_name ] } ] } ] }
```



MY IP OR YOUR IP?


- We can reference an existing virtual network
- Or specify a fixed IP address



```
vip_fixed:  
  type: OS::ContrailV2::InstanceIp  
  depends_on: [ vmi_obj ]  
  properties:  
    virtual_machine_interface_refs: [{ get_resource: vmi_obj }]  
    virtual_network_refs: [{ get_param: mgmt-net }]  
    instance_ip_address: '192.168.1.4'
```

VIP ADDRESS

- Many VMs sharing an IP
- To spread the load on multiple devices
- Leveraging Contrail ECMP capabilities



ALLOWING MULTIPLE ADDRESSES ON A PORT



192.168.1.13

192.168.1.14

NO



YES



192.168.1.14

192.168.1.14

VIP MEANS ECMP

- VIP configured on the VMI
- Need to allow an Allowed Address Pair
- Routes automatically added in Contrail
- Multiple routes towards the same VIP are seen as ECMP paths

Prefix: 192.168.5.100 / 32 (4 Routes)									
▶ ECMP Composite sub nh count: 2	Source IP:	Destination IP:	vrf:	Ref count:	Policy: enabled	Peer: 192.168.100.2	Valid: true	Label: -1	
▶ ECMP Composite sub nh count: 2	Source IP:	Destination IP:	vrf:	Ref count:	Policy: enabled	Peer: Ecmp	Valid: true	Label: 20	
▶ interface	Interface: tap24c1f7ca-fe	Destination VN: default-domain:admin:vmi-net	Policy: enabled	Peer: LocalVmPort	Valid: true				
▶ interface	Interface: tapcfed0090-ec	Destination VN: default-domain:admin:vmi-net	Policy: enabled	Peer: LocalVmPort	Valid: true				


ALLOWED ADDRESS PAIR

```
resources:
  vmi_1:
    type: OS::ContrailV2::virtualMachineInterface
    properties:
      name: 'my_vmi'
      virtual_network_refs: [{ list_join: [':', { get_attr: [ new_vn, fq_name ] } ] }]
      virtual_machine_interface_mac_addresses:
        {
          virtual_machine_interface_mac_addresses_mac_address: 'aa:bb:cc:dd:ee:ff',
        }
      virtual_machine_interface_allowed_address_pairs:
        {
          virtual_machine_interface_allowed_address_pairs_allowed_address_pair: [{
            virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip:
              {
                virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip_ip_prefix: '192.168.1.100',
                virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip_ip_prefix_len: '32',
              },
            virtual_machine_interface_allowed_address_pairs_allowed_address_pair_mac: 'aa:bb:cc:dd:ee:ff',
            virtual_machine_interface_allowed_address_pairs_allowed_address_pair_address_mode: 'active-active',
          }]
        }
      }
```



MULTIPLE IPS ON THE SAME MAC

- Configuration within the VMI object
- Add allowed pairs
 - MAC address
 - IP address



CONTRAIL PICKS THE MAC

- We can avoid to define a custom MAC
- Just remove MAC definition and AAP MAC reference

```
resources:
  vmi_1:
    type: OS::ContrailV2::VirtualMachineInterface
    properties:
      name: 'my_vmi'
      virtual_network_refs: [{ list_join: [':', { get_attr: [ new_vn, fq_name ] } ] }]
      virtual_machine_interface_allowed_address_pairs:
        {
          virtual_machine_interface_allowed_address_pairs_allowed_address_pair: [{
            virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip:
              {
                virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip_ip_prefix: { get_param: '192.168.1.100' },
                virtual_machine_interface_allowed_address_pairs_allowed_address_pair_ip_ip_prefix_len: { get_param: '32' },
              },
            virtual_machine_interface_allowed_address_pairs_allowed_address_pair_address_mode: { get_param: 'active-active' },
          }]
        }
      }
```



VIP IN ACTION

```
Connected (unencrypted) to: QEMU (instance-00000064)
$ ssh cirros@192.168.5.100


Host '192.168.5.100' is not in the trusted hosts file.
(ecdsa-sha2-nistp521 fingerprint md5 f7:7b:6a:43:3a:fa:f5:3f:2a:ac:5d:db:4c:10:7c:b3)
Do you want to continue connecting? (y/n) ^C
$ ssh cirros@192.168.5.100

Host '192.168.5.100' is not in the trusted hosts file.
(ecdsa-sha2-nistp521 fingerprint md5 bc:24:84:5a:91:f0:3c:81:96:9e:83:5a:fa:57:a0:e1)
Do you want to continue connecting? (y/n) ^C
$ ssh cirros@192.168.5.100

Host '192.168.5.100' is not in the trusted hosts file.
(ecdsa-sha2-nistp521 fingerprint md5 f7:7b:6a:43:3a:fa:f5:3f:2a:ac:5d:db:4c:10:7c:b3)
Do you want to continue connecting? (y/n) ^C
```


TRUNK PORTS

- Multiple vlans inside the VM
- Virtual NIC must become a Trunk
- Do not Use Openstack Trunk Object
- Use Contrail sub-interfaces (which came earlier!)



HOW DO TRUNKS WORK?

- Concept is easy
- Create a parent network
- Attach a VMI to it
- Create child virtual networks, one for each vlan
- Create child VMIs, one for each vlans
- Child VMIs must reference the parent VMI
- Child VMIs must use parent VMI MAC address



FIRST THERE IS A FATHER

```
resources:
  parent_vn:
    type: OS::ContrailV2::virtualNetwork
    properties:
      name: { get_param: parent }
      network_ipam_refs: ['default-domain:my-project:default-ipam']
      network_ipam_refs_data:
        [{
          network_ipam_refs_data_ipam_subnets:
            [{
              network_ipam_refs_data_ipam_subnets_subnet:
                {
                  network_ipam_refs_data_ipam_subnets_subnet_ip_prefix: '172.30.1.0',
                  network_ipam_refs_data_ipam_subnets_subnet_ip_prefix_len: '24',
                },
            ]
        }]
  vm_1:
    type: OS::Nova::Server
    properties:
      name: 'trunk_vm'
      image: 'vsrx'
      flavor: 'vsrx'
      networks:
        - network: { get_resource: parent_vn }
```



parent VMI

THEN THE CHILDREN

- Child VN is a standard VN
- No vlan defined here
- No link to parent


```
child_vn_1:
  type: OS::ContrailV2::VirtualNetwork
  properties:
    name: 'child1'
    network_ipam_refs: ['default-domain:my-project:default-ipam']
    network_ipam_refs_data:
      [{
        network_ipam_refs_data_ipam_subnets:
          [{
            network_ipam_refs_data_ipam_subnets_subnet:
              {
                network_ipam_refs_data_ipam_subnets_subnet_ip_prefix: '172.30.100.0',
                network_ipam_refs_data_ipam_subnets_subnet_ip_prefix_len: '24',
              },
          ]
      }]
  ]
```



MULTIPLE PORTS, ONE MAC

- Child VMI is where the magic happens
- We reference another VMI, the parent
- We set the MAC to parent MAC VMI address
- We assign the vlan

```
vmi_1:
  type: OS::ContrailV2::VirtualMachineInterface
  properties:
    name: 'vmi_1'
    virtual_network_refs: [{ get_resource: child_vn_1 }]
    virtual_machine_interface_properties:
      {
        virtual_machine_interface_properties_sub_interface_vlan_tag: '100',
      }
    virtual_machine_interface_mac_addresses:
      {
        virtual_machine_interface_mac_addresses_mac_address: [{ get_attr: [vm_1, addresses, get_param: 'parent', 0,
"OS-EXT-IPS-MAC:mac_addr"] }],
      }
    virtual_machine_interface_refs: [{ get_attr: [vm_1, addresses, get_param: 'parent', 0, "port"] }]
```



HOW DO I GET THE RIGHT MAC?

- We need to obtain the MAC of the parent
- By checking VM attributes
- We extract it from ports attributes
- Only check on the port connected to the parent network



```
virtual_machine_interface_mac_addresses:  
  {  
    virtual_machine_interface_mac_addresses_mac_address: [{ get_attr: [vm_1,  
addresses, get_param: 'parent', 0, "OS-EXT-IPS-MAC:mac_addr"] }],  
  }
```


SAME OLD INSTANCE IP


- Still need the Instance IP
- But nothing new here
- Same old Instance IP object



```
iip_1:
  type: OS::ContrailV2::InstanceIp
  depends_on: [ vmi_1 ]
  properties:
    virtual_network_refs: [{ get_resource: child_vn_1 }]
    virtual_machine_interface_refs: [{ get_resource: vmi_1 }]
```

GOING OUT THE DC

- My Virtual Network needs to leave the DC
- In Openstack we have Provider Networks
- Here still tenant networks
- But we use Route targets
- Then same old MP-BGP VPN mechanisms between Contrail and SDN gateway



SDN GATEWAY AS A PE

- Need configuration on SDN gateway as well
- We need a VRF
- Matching route-target
- Usual VRF import/export policies

```
root@sdn-gw# show routing-instances test_vn
instance-type vrf;
interface lo0.1000;
route-distinguisher 192.168.100.1:1;
vrf-import test_imp;
vrf-export test_exp;
vrf-table-label;
```



THE RIGHT TARGET

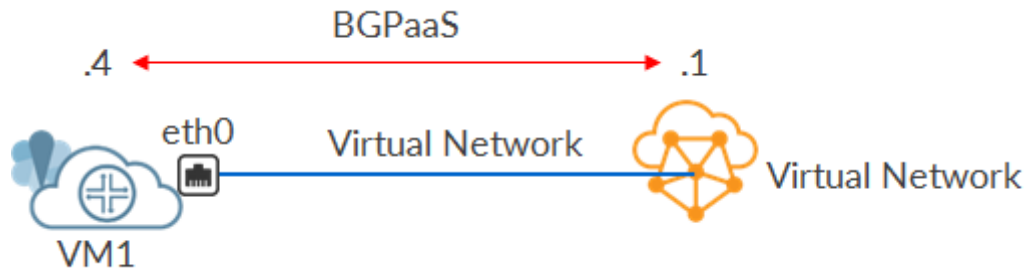
- Virtual Network can have its own route target
- Can import other route-targets
- Can export route-targets

```
resources:
  vip_vn:
    type: OS::ContrailV2::VirtualNetwork
    properties:
      name: { get_param: name}
      import_route_target_list:
        {
          import_route_target_list_route_target: [ 'target:65500:131' , 'target:65500:241' ]
        }
      export_route_target_list:
        {
          export_route_target_list_route_target: [ 'target:65500:201' ]
        }
      route_target_list:
        {
          route_target_list_route_target: ['target:65500:101'],
        }
```



BGP AS A SERVICE

- Totally new feature
- VMs by default are blind to vRouter networking information
- Need to create a communication between vRouter and the VM
- This is BGP as a Service (BGPaaS)



DO YOU TALK BGP?

- VM must talk BGP
- Session with Virtual Network gateway address
- Additional options like prefix limit, as-override, etc...
- BGPaaS object references a VMI





CUSTOMIZE YOUR BGPAAS

```
bgpaas_2:
  type: OS::ContrailV2::BgpAsService
  properties:
    name: 'bgp-session-1'
    virtual_machine_interface_refs: [{ get_resource: vmi_2  }]
    autonomous_system: '65511'
    bgpaas_session_attributes:
      {
        bgpaas_session_attributes_address_families:
          {
            bgpaas_session_attributes_address_families_family: ['inet','inet6'],
          },
        bgpaas_session_attributes_family_attributes:
          [{
            bgpaas_session_attributes_family_attributes_address_family: inet,
            bgpaas_session_attributes_family_attributes_loop_count: 1,
            bgpaas_session_attributes_family_attributes_prefix_limit:
              {
                bgpaas_session_attributes_family_attributes_prefix_limit_maximum: 10
              }
          }]
        bgpaas_session_attributes_as_override: 'true'
        bgpaas_session_attributes_hold_time: '30'
        bgpaas_session_attributes_private_as_action: 'remove'
      }
  }
```




NEUTRON VS CONTRAIL ?

- Neutron objects are available
- We can use them with Contrail
- Not all Contrail objects are compatible with Contrail



WHEN DO WE HAVE TO CHANGE?

- You need additional features that only Contrail can offer
- For example, ECMP fields
- You need to implement something that, conceptually, is totally different between Contrail and vanilla Neutron
- For example, VN that need to go outside the DC



LOST IN TRANSLATION




- We start from a Neutron based template
- And need to move to a Contrail based one
- We tried to identify some common use-cases
- And give best practices on how to transform a stack of Neutron objects in one made with Contrail objects



CREATING SERVICES IS LIKE A PUZZLE!




BASIC TRANSLATION




TRANSLATION #1

- Create a VM with one vNIC connected to a new virtual network



HEAT MAP #1




TRANSLATION #2

- Create a VM with one vNIC connected to an existing virtual network



HEAT MAP #2

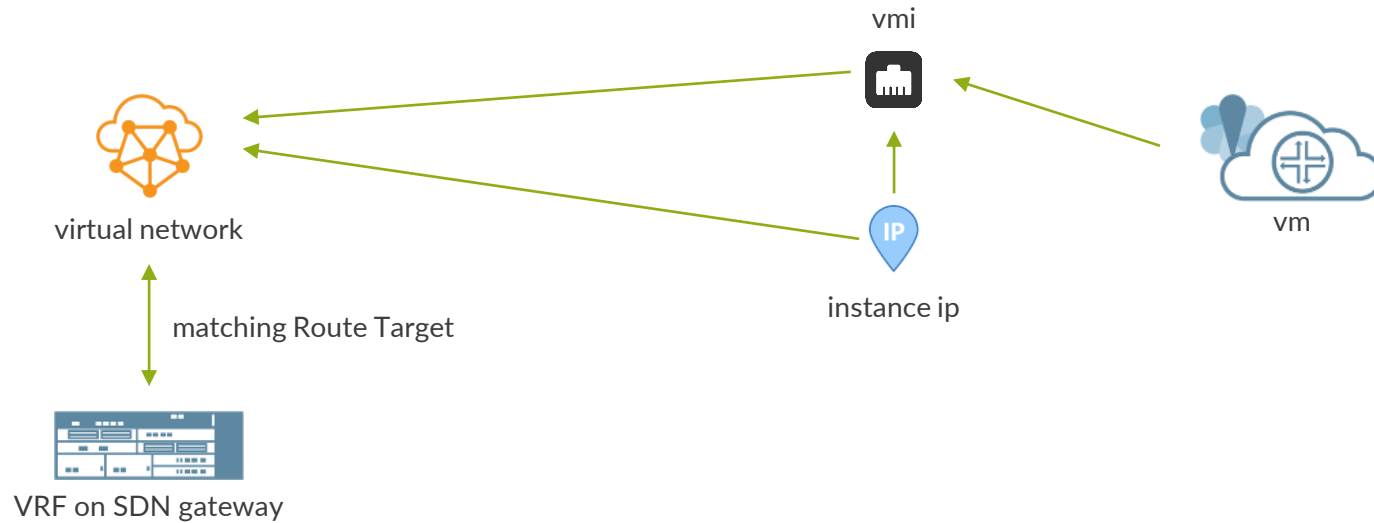


TRANSLATION #3

- Create a VM with one vNIC connected to a new virtual network that needs to go outside the DC



HEAT MAP #3

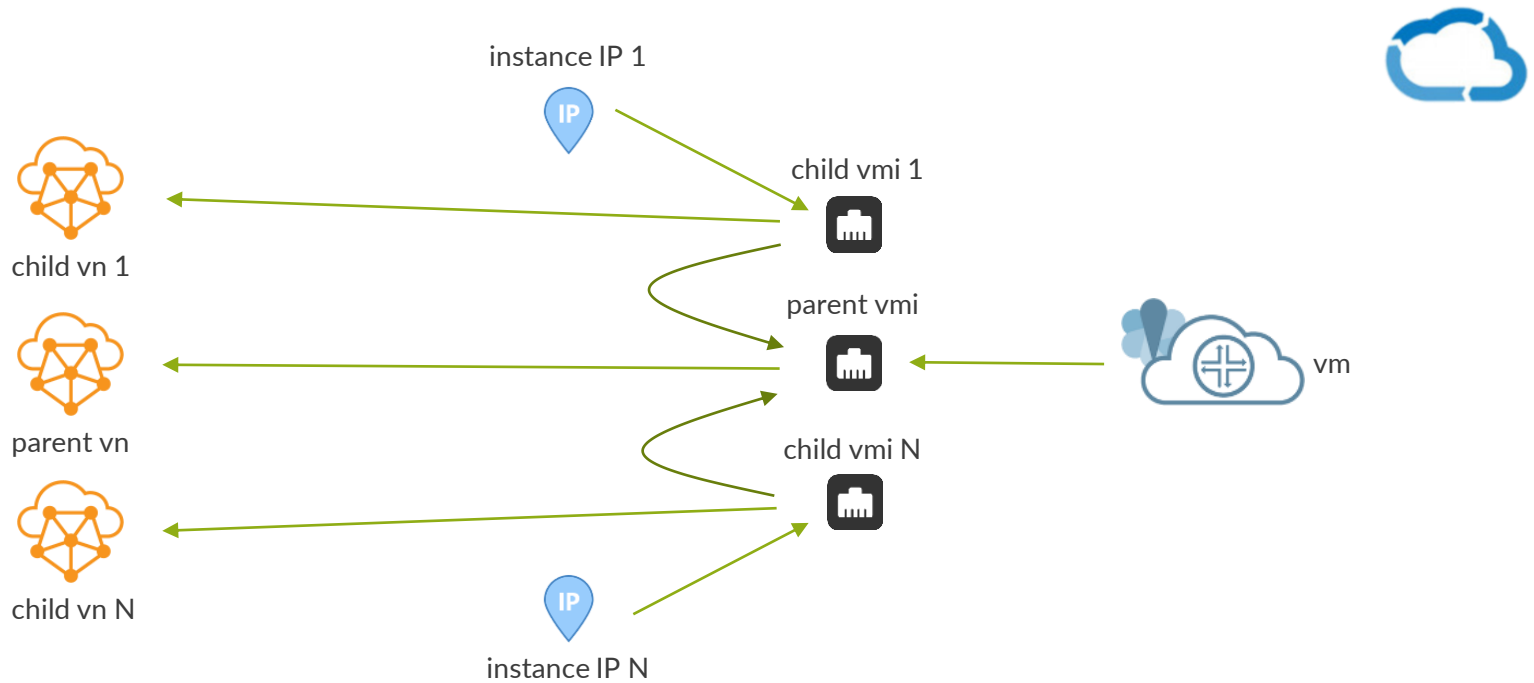


TRANSLATION #4

- The VM vNIC needs to be configured as trunk with N vlans




HEAT MAP #5




TRANSLATION #5

- The VM needs to be connected to a L2 network, VXLAN ready



HEAT MAP #5



TRANSLATION #6

- The virtual network needs to use a custom gw and a custom allocation pool



Neutron Net



Neutron Subnet




Neutron Subnet pools



Virtual Network
with custom gateway
& Allocation pools

HEAT MAP #6




TRANSLATION #7

- The VM must have a vNIC with a custom MAC, a fixed IP and must allow the configuration of a VIP address



HEAT MAP #7

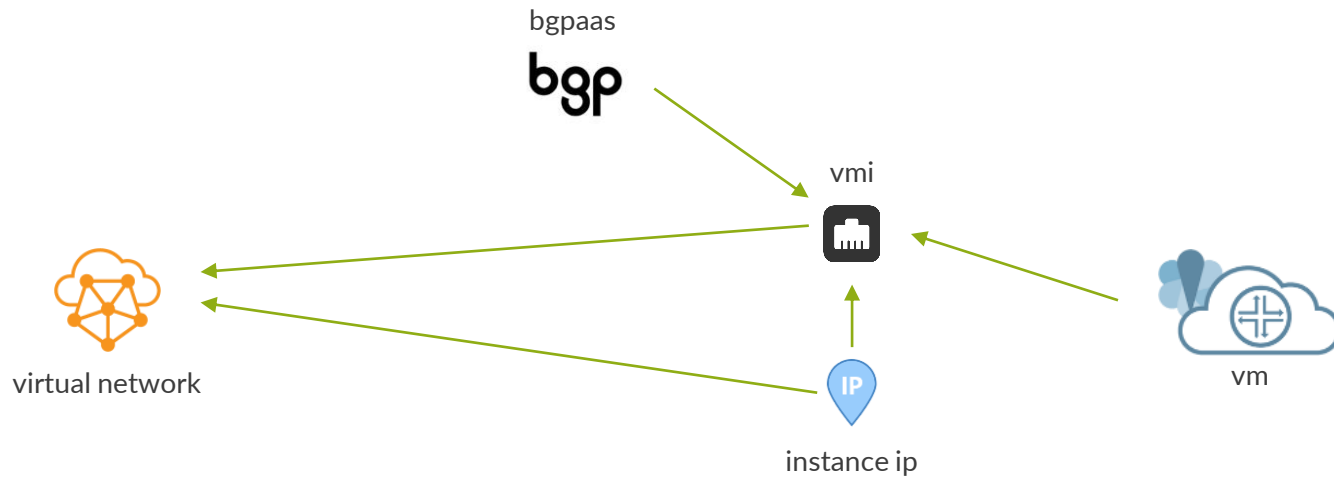


TRANSLATION #8

- The VM must talk BGP in order to exchange BGP routers and use BFD for faster failure detection

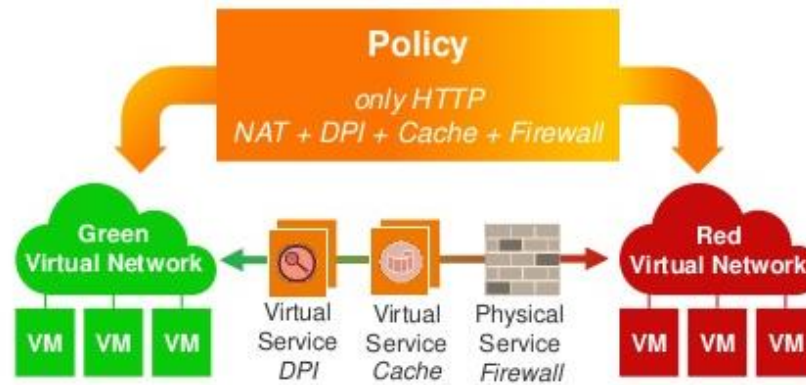


HEAT MAP #8



SERVICE CHAINING

- Contrail exclusive feature
- Manage communications between virtual networks
- By adding services between them!




AT FIRST IT WAS TWO NETWORKS

- By default two virtual networks cannot talk to each other
- They are separate broadcast domain
- In a physical network we would need a router




THEN A POLICY CONNECTED THEM

- Networks can talk to each other through a network policy
- Applied to both virtual networks
- A list of rules telling which traffic can pass or not
- Rules built over protocol, dst/src addresses, dst/src virtual networks, dst/src ports



AND A SERVICE CONTROLLED COMMUNICATIONS

- If policy allows traffic to pass we can add one or more service between the two virtual networks
- Service can be a firewall
- We bring security between two virtual networks
- Service chaining routing is automatically managed by Contrail



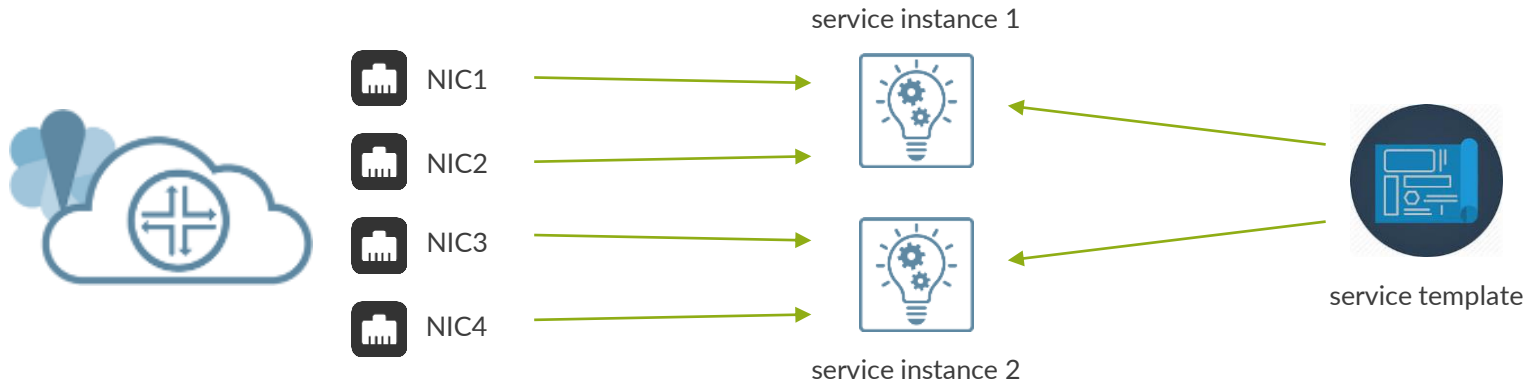
SERVICE CHAIN RECIPE

- Two virtual networks
- One network policy applied to those networks
- One or more VMs acting as service elements
- Service templates describing the nature of the service instance
- Port tuple, a collection of ports belonging to a VM
- Service instance, the actual service part of the service chain



DON'T BE CONFUSED!

- Based on the template
- Uses ports of a VM as defined in a port tuple
- A service instance is not a VM!
- We might have multiple service instances using the same VM



SERVICE CHAIN BUILDING BLOCKS

- Network policy
- Service template
- Service instance
- Port tuple
- Modified VMI
- Service VM



WRITING YOUR POLICY

```
type: OS::ContrailV2::NetworkPolicy
properties:
  name: { get_param: pol_name }
  network_policy_entries: {
    network_policy_entries_policy_rule: [{
      network_policy_entries_policy_rule_direction: { get_param: direction },
      network_policy_entries_policy_rule_protocol: { get_param: protocol },
      network_policy_entries_policy_rule_src_ports: [{
        network_policy_entries_policy_rule_src_ports_start_port: { get_param: src_port_start },
        network_policy_entries_policy_rule_src_ports_end_port: { get_param: src_port_end }
      }],
      network_policy_entries_policy_rule_dst_ports: [{
        network_policy_entries_policy_rule_dst_ports_start_port: { get_param: dst_port_start },
        network_policy_entries_policy_rule_dst_ports_end_port: { get_param: dst_port_end }
      }],
      network_policy_entries_policy_rule_dst_addresses: [{
        network_policy_entries_policy_rule_dst_addresses_virtual_network: { get_param: dst_add }
      }],
      network_policy_entries_policy_rule_src_addresses: [{
        network_policy_entries_policy_rule_src_addresses_subnet: {
          network_policy_entries_policy_rule_src_addresses_subnet_ip_prefix: {get_param: rule_cidr_1},
          network_policy_entries_policy_rule_src_addresses_subnet_ip_prefix_len: {get_param: rule_mask_1},
        }
      }],
      network_policy_entries_policy_rule_action_list: {
        network_policy_entries_policy_rule_action_list_simple_action: { get_param: simple_action },
        network_policy_entries_policy_rule_action_list_apply_service: [{ get_param: svc_inst_freq_name }]
      },
    }
  ]
}
```



APPLYING POLICY TO A VIRTUAL NETWORK

- Policy must be assigned to virtual networks too
- To both virtual networks, left and right

```
left_vn:
  type: OS::ContrailV2::VirtualNetwork
  depends_on: [ net_policy_apn_1 ]
  properties:
    ...
    network_policy_refs: [{ get_param: pol_fq_name }]
    network_policy_refs_data:
      [{
        network_policy_refs_data_sequence:
          {
            network_policy_refs_data_sequence_major: 0,
            network_policy_refs_data_sequence_minor: 0,
          },
      }]
  ]
```

DRAW YOUR TEMPLATE

- Defines service characteristics
- Type, mode, interfaces

```
svc_tmplt:
  type: OS::ContrailV2::ServiceTemplate
  properties:
    name: { get_param: service_template_name }
    service_template_properties:
      {
        service_template_properties_version: '2',
        service_template_properties_service_mode: 'in-network',
        service_template_properties_service_type: 'firewall',
        service_template_properties_interface_type:
          [
            {
              service_template_properties_interface_type_service_interface_type: 'left',
            },
            {
              service_template_properties_interface_type_service_interface_type: 'right',
            }
          ],
        service_template_properties_ordered_interfaces: 'true',
      }
    domain: 'default-domain'
```



DEFINE THE SERVICE INSTANCE

- We bind interfaces to virtual networks

```
svc_inst_apn_1:
  type: OS::Contrailv2::ServiceInstance
  depends_on: [ svc_tmplt ]
  properties:
    name: { get_param: svc_inst_name }
    service_instance_properties:
      {
        service_instance_properties_interface_list:
          [
            {
              service_instance_properties_interface_list_virtual_network:
                {
                  list_join: [':', { get_attr: [ left_vn, fq_name ] } ]
                },
            },
            {
              service_instance_properties_interface_list_virtual_network:
                {
                  list_join: [':', { get_attr: [ right_vn, fq_name ] } ]
                },
            },
          ],
        }
    service_template_refs: [{ get_resource: svc_tmplt }]
```




GATHER PORTS INTO A TUPLE

- Binds a list of interfaces
- References a service instance

```
prt_tpl_1:
  type: OS::ContrailV2::PortTuple
  depends_on: [ svc_inst_apn_1 ]
  properties:
    name: { get_param: tuple_name }
    service_instance: { list_join: [':', { get_attr: [ svc_inst_apn_1, fq_name ] } ] }


vmi_2:
  type: OS::ContrailV2::VirtualMachineInterface
  depends_on: [ prt_tpl_1 ]
  properties:
    virtual_machine_interface_properties:
      {
        virtual_machine_interface_properties_service_interface_type: 'left',
      }
    port_tuple_refs: [{ get_resource: prt_tpl_1 }]
    virtual_network_refs: [{ list_join: [':', { get_attr: [ left_vn, fq_name ] } ] }]
```



SERVICE VM WILL MANAGE TRAFFIC

- Standard VM
- Some of its ports assigned to the service instance

```
vm1:
  type: OS::Nova::Server
  depends_on: [ iip_1, iip_2, iip_3 ]
  properties:
    name: { get_param: vm1name }
    image: { get_param: image }
    flavor: { get_param: flavor }
    networks:
      - port: { get_resource: vmi_1 }
      - port: { get_resource: vmi_2 }
      - port: { get_resource: vmi_3 }
```



SERVICE CHAINING HEAT MAP



left vn



right vn



policy



service instance




port tuple



service template

VM

- x3 vmi
- x3 instance IP



USE CASES? SKY IS THE LIMIT

- Policy allows inter-VN communication
- Service configured between the two VN
- Service Firewall blocks some traffic!



SUMMING UP

- SDN brings automation
- Automation in Openstack means Heat
- Contrail has its own Heat objects
- Models to create VNFs with Contrail
- Best practices to translate Neutron templates to Contrail templates

SUMMARY



THAT'S ALL!

